

Welcome to

CS 3516:
Computer Networks

Prof. Yanhua Li

Time: 9:00am –9:50am M, T, R, and F

Location: AK219

Fall 2018 A-term

Quiz 5 on Friday

P2P networks

Mid-term next Friday (Tentative)

Lab 2 due on Friday

Piazza is available

Chapter 2: outline

2.6 P2P applications

1. P2P vs Client&Server


2. Unstructured Peer-to-Peer Networks

BitTorrent

3. Structured Peer-to-Peer Networks

Distributed Hash Table (DHT)

Client-server vs P2P architecture



Peer-to-Peer Networks:

❖ Unstructured Peer-to-Peer Networks

- Napster
- Gnutella
- BitTorrent

Chapter 2: outline

2.6 P2P applications

1. P2P vs Client&Server


2. Unstructured Peer-to-Peer Networks

- Napster
- Gnutella
- BitTorrent

3. Structured Peer-to-Peer Networks

Peer-to-Peer Networks: How Did it Start?

- ❖ A killer application: Napster
 - Free music over the Internet
- ❖ Key idea: **share the content, storage *and* bandwidth of individual (home) users**



Model


- ❖ Each user stores a subset of files
- ❖ Each user has access (can download) files from all users in the system

Challenges


- ❖ Scale: up to hundred of thousands or millions of machines
- ❖ Dynamicity: machines can come and go any time

Main Challenge

- ❖ Find where a particular file is stored



Napster: Example



Peer-to-Peer Networks: Napster

❖ Napster history: the rise

- January 1999: Napster version 1.0
- May 1999: company founded
- September 1999: first lawsuits
- 2000: 80 million users



**Shawn Fanning,
Northeastern freshman**


❖ Napster history: the fall

- Mid 2001: out of business due to lawsuits
- Mid 2001: dozens of P2P alternatives that were harder to touch, though these have gradually been constrained
- 2003: growth of pay services like iTunes

❖ Napster history: the resurrection

- 2003: Napster reconstituted as a pay service
- 2006: still lots of file sharing going on

Napster: Example



- ❖ User installing the software
- ❖ Client contacts Napster (via TCP)
- ❖ Client searches on a title or performer
- ❖ Client requests the file from the chosen supplier

Napster: Limitations of Central Directory

- ❖ Single point of failure
- ❖ Performance bottleneck
- ❖ Copyright infringement

File transfer is decentralized, but locating content is highly centralized

- ❖ So, later P2P systems were more distributed

Peer-to-Peer Networks: Gnutella

❖ Gnutella history

- 2000: J. Frankel & T. Pepper **released Gnutella**
- **Soon after:** many other clients (e.g., Morpheus, Limewire, Bearshare)
- **2001:** protocol enhancements, e.g., “ultrapeers”

❖ Query flooding


- **Join:** contact a few nodes to become neighbors
- **Publish:** no need!
- **Search:** ask neighbors, who ask their neighbors
- **Fetch:** get file directly from another node



gnutella.com


Gnutella

- ❖ Ad-hoc topology
- ❖ No guarantees on recall




- ❖ Queries are flooded for bounded number of hops (TTL)

Gnutella: Protocol

- Query message sent over existing TCP connections
 - Peers forward Query message
 - QueryHit sent over reverse path
- 

Scalability:
limited scope
flooding



Edge between peer X and Y if there's a TCP connection

Gnutella: Pros and Cons

❖ Advantages

- Fully decentralized, Highly robust

❖ Disadvantages

- Not scalable; the entire network can be swamped with request
 - Search scope may be quite large, High overhead
 - Search time may be quite long

Chapter 2: outline

2.6 P2P applications

1. P2P vs Client&Server
2. Unstructured Peer-to-Peer Networks
BitTorrent

P2P design challenges

- ❖ Scalability

- Publish/Download

- ❖ Large file

- ❖ Free Riding

BitTorrent: Simultaneous Downloading

- ❖ Divide large file into many pieces (256Kbytes)
 - Replicate different pieces on different peers
 - A peer with a complete piece can trade with other peers
 - Peer can (hopefully) assemble the entire file
- ❖ Allows simultaneous downloading
 - Retrieving different parts of the file from different peers at the same time

BitTorrent Components

❖ Seed

- Peer with entire file
- Fragmented in pieces

❖ Leech

- Peer with an incomplete copy of the file


❖ Torrent file

- Passive component
- Stores summaries of the pieces to allow peers to verify their integrity


❖ Tracker

- Allows peers to find each other
- Returns a list of random peers

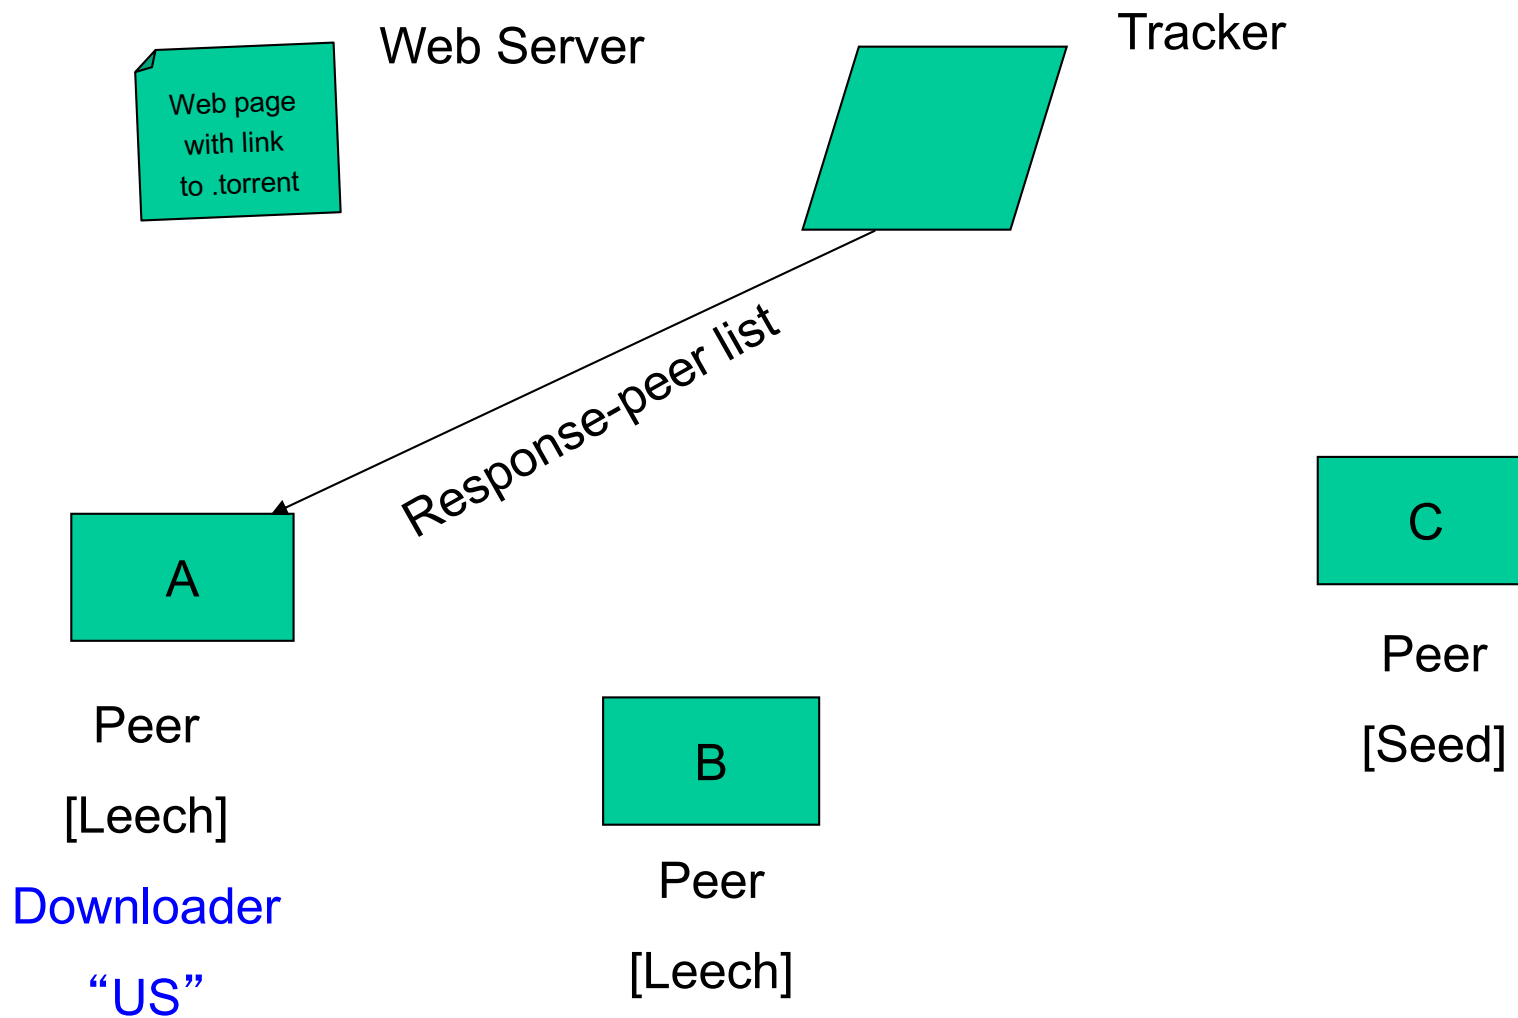
BitTorrent: Overall Architecture



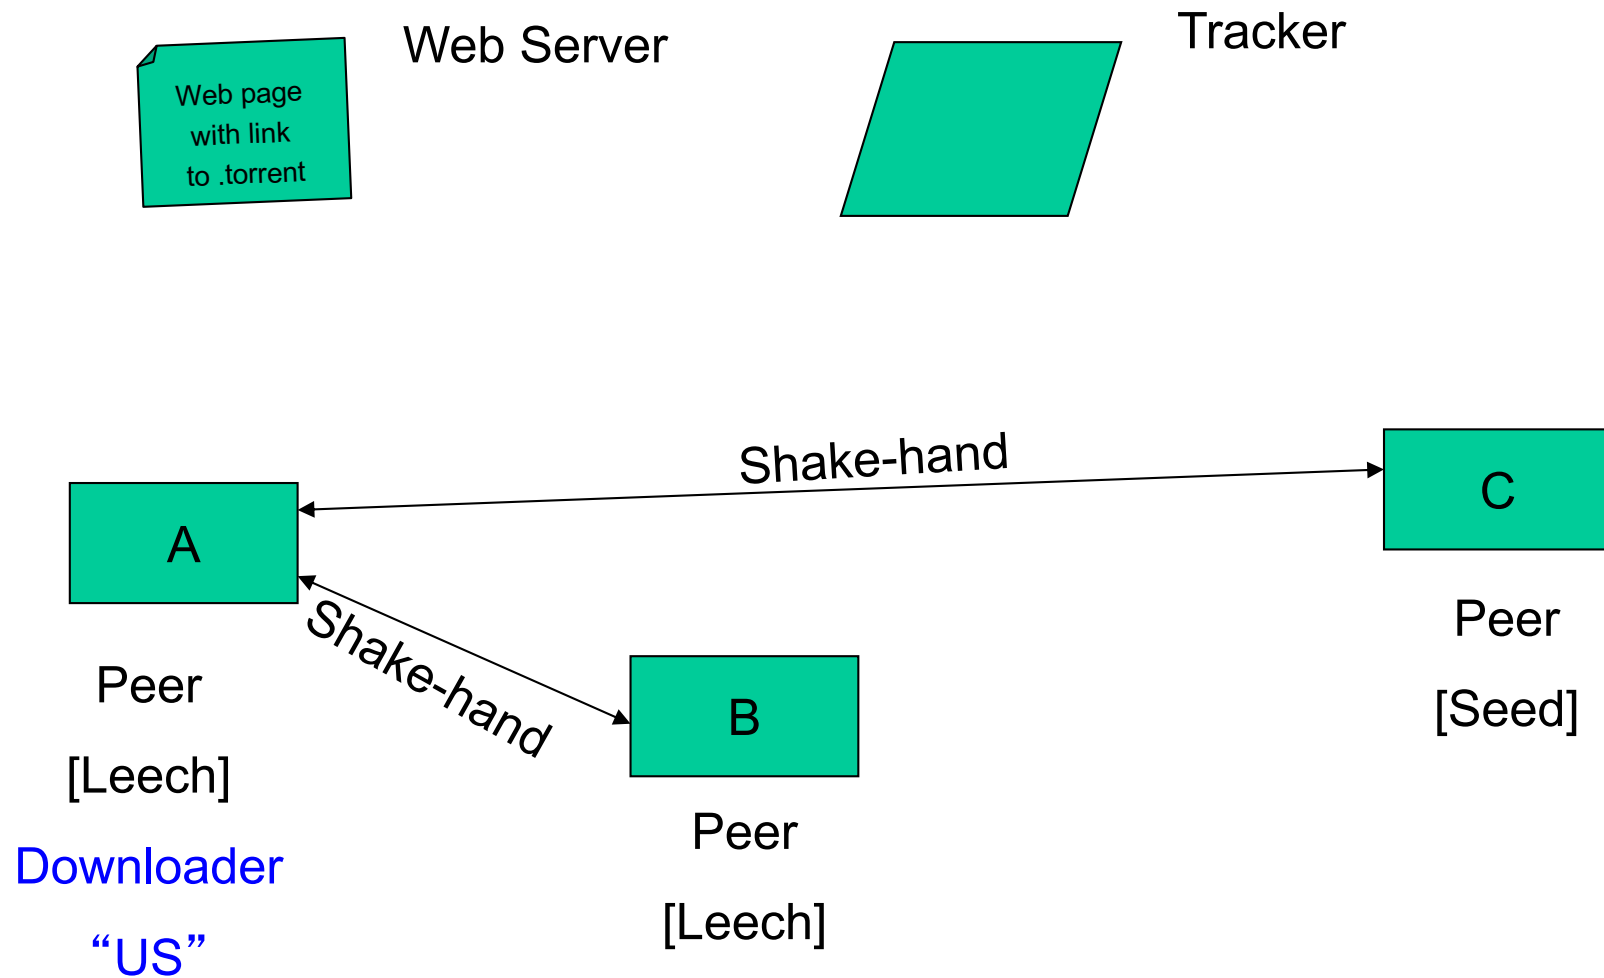
BitTorrent: Overall Architecture



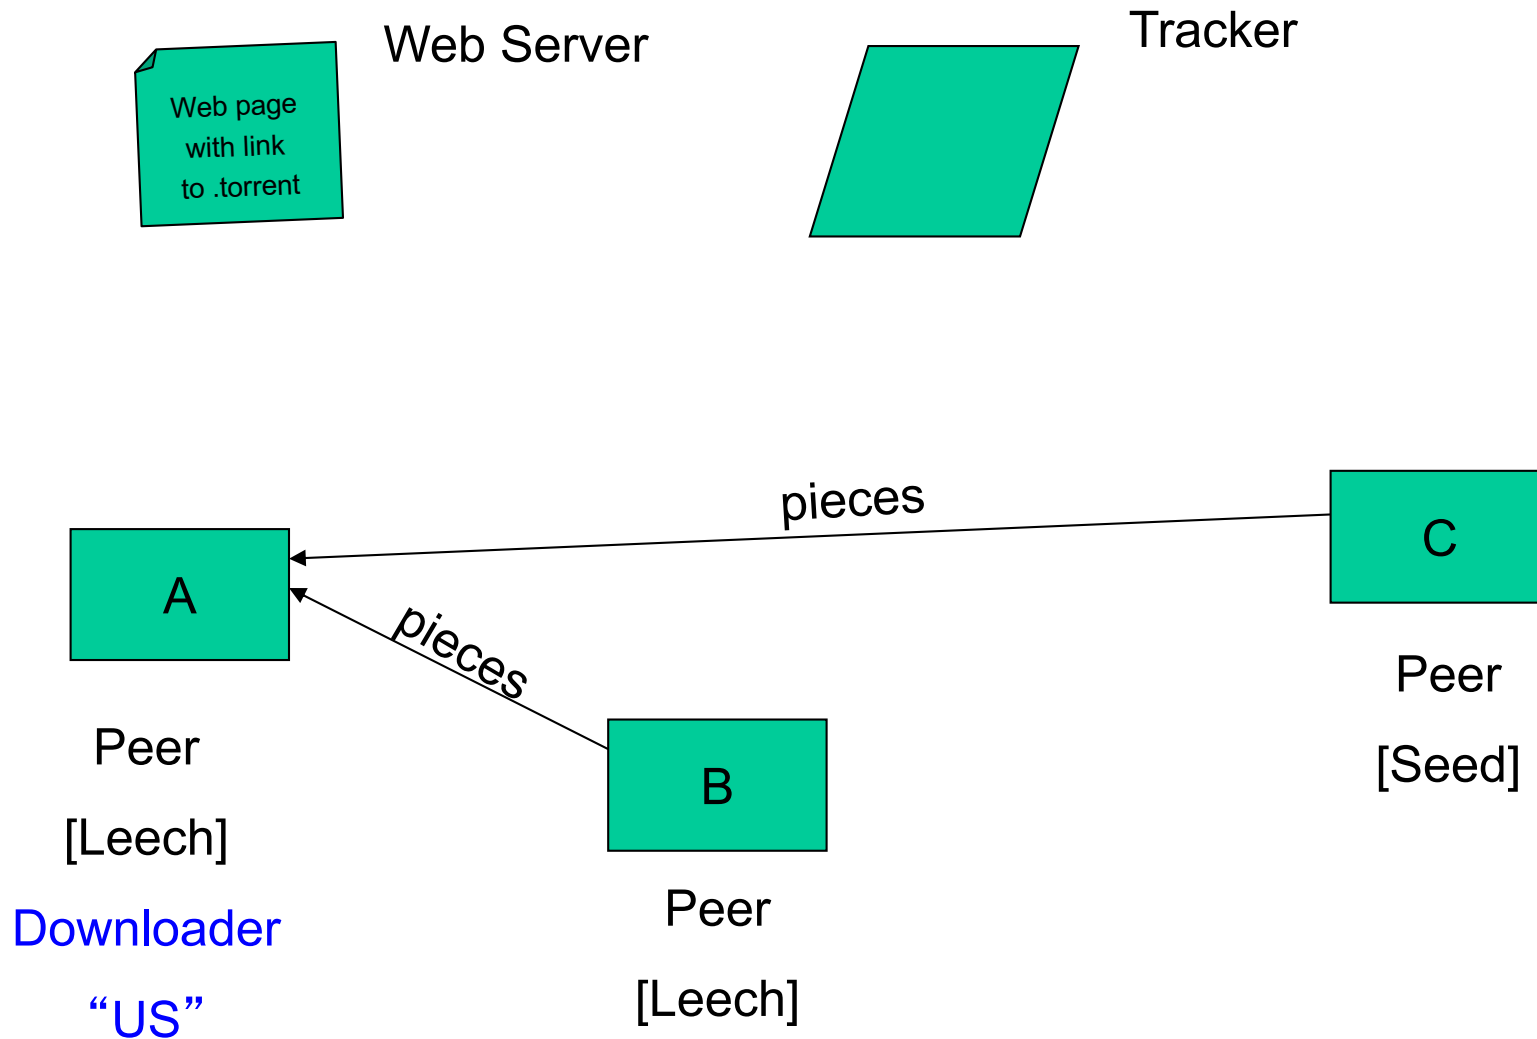
BitTorrent: Overall Architecture




BitTorrent: Overall Architecture




BitTorrent: Overall Architecture



BitTorrent: Overall Architecture



BitTorrent: Overall Architecture



Free-Riding Problem in P2P Networks

- ❖ Vast majority of users are free-riders
 - Most share no files and answer no queries
 - Others limit # of connections or upload speed
- ❖ A few “peers” essentially act as servers
 - A few individuals contributing to the public good
 - Making them hubs that basically act as a server
- ❖ BitTorrent prevent free riding
 - Allow the fastest peers to download from you
 - Occasionally let some free loaders download

BitTorrent: requesting, sending file chunks

requesting chunks:

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ *periodically*, Alice asks each peer for list of chunks that they have
- ❖ Alice requests missing chunks from peers, *rarest first*


sending chunks: tit-for-tat

- ❖ Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - *re-evaluate top 4 every 10 secs*
- ❖ *every 30 secs: randomly select another peer, starts sending chunks*
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4



BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Chapter 2: summary

our study of network apps now complete!

- ❖ application architectures
 - client-server
 - P2P
- ❖ application service requirements:
 - reliability, throughput, delay, security
- ❖ Internet transport service model
 - connection-oriented, reliable: TCP
 - unreliable, datagrams: UDP
- ❖ specific protocols:
 - HTTP
 - SMTP, POP, IMAP
 - DNS
 - P2P: BitTorrent

Questions?