

Welcome to

***CS 3516:***  
***Computer Networks***

Prof. Yanhua Li

Time: 9:00am –9:50am M, T, R, and F

Location: AK219

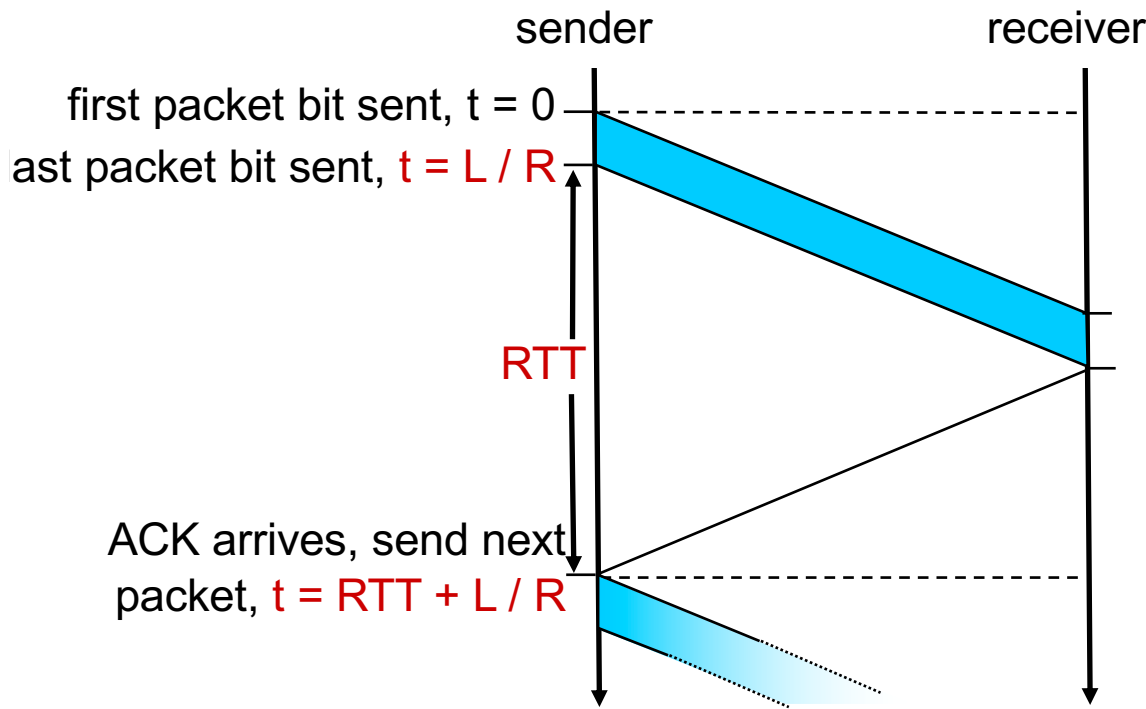
Fall 2018 A-term

# Updates

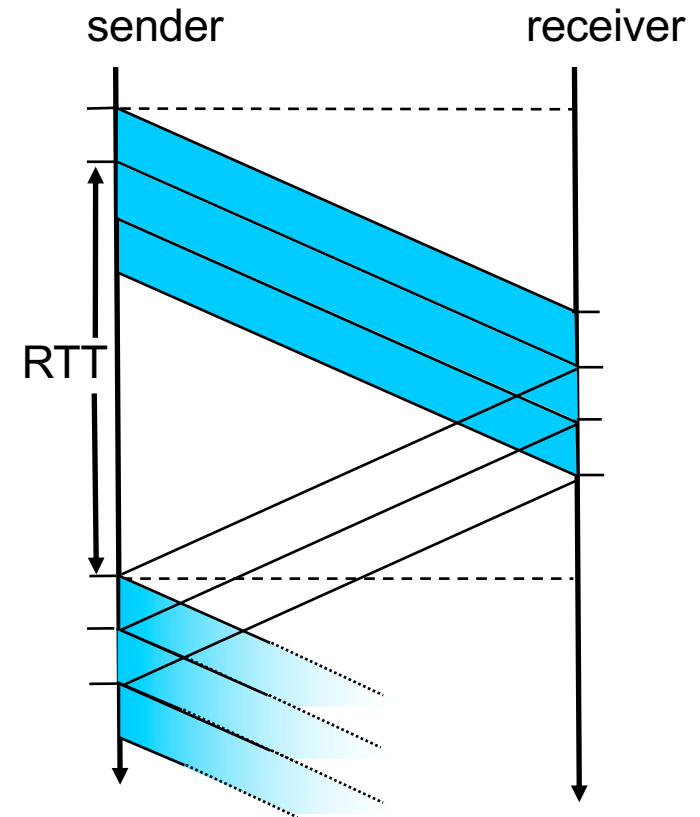
---

- ❖ Quiz 6
  - Grading by Today
- ❖ Mid-term
  - Grading by Wed
- ❖ Project 2
  - Due on 9/28 F
  - Extra office hours (TBD)
- ❖ Quiz 7:
  - This Thursday
  - *TCP and Network Layer Intro*

# rdt3.0 vs Pipelining approach



RDT3.0



Pipelining approach

3-packet pipelining increases utilization by a factor of 3!

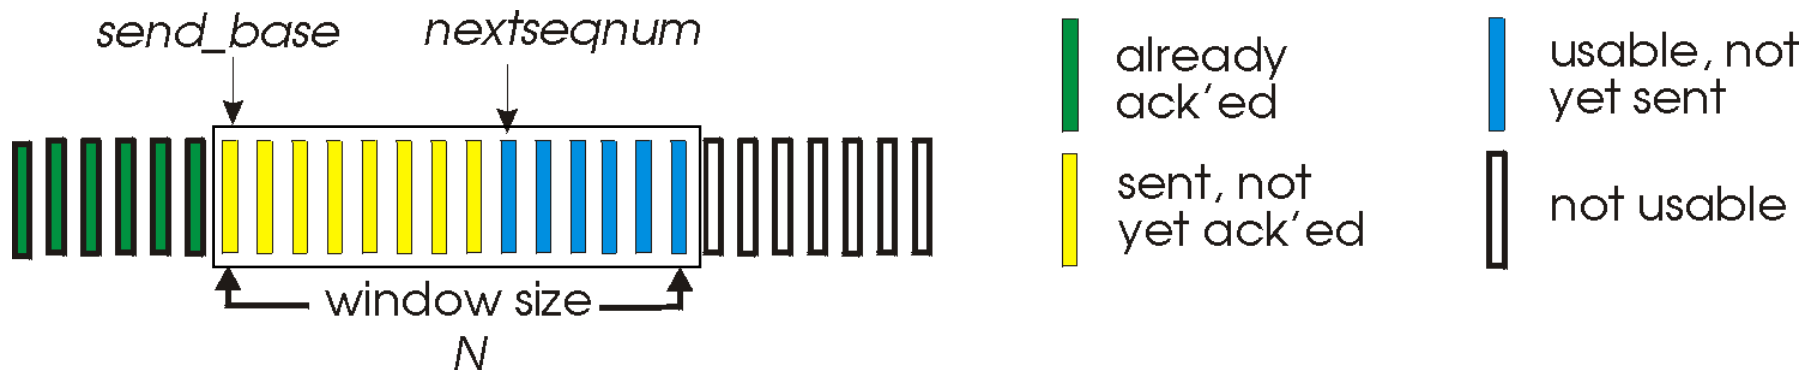
# Pipelined protocols: overview

## Go-back-N:

- sender can have up to N unacked packets in pipeline
  - receiver only sends *cumulative ack*
    - doesn't ack packet if there's a gap
  - sender has timer for oldest unacked packet
    - when timer expires, retransmit *all* unacked packets
- ❖ Extending from
  - ❖ **one** unacknowledged pkt (in RDT3.0) to
  - ❖ **multiple** unacknowledged pkts (in pipelining)

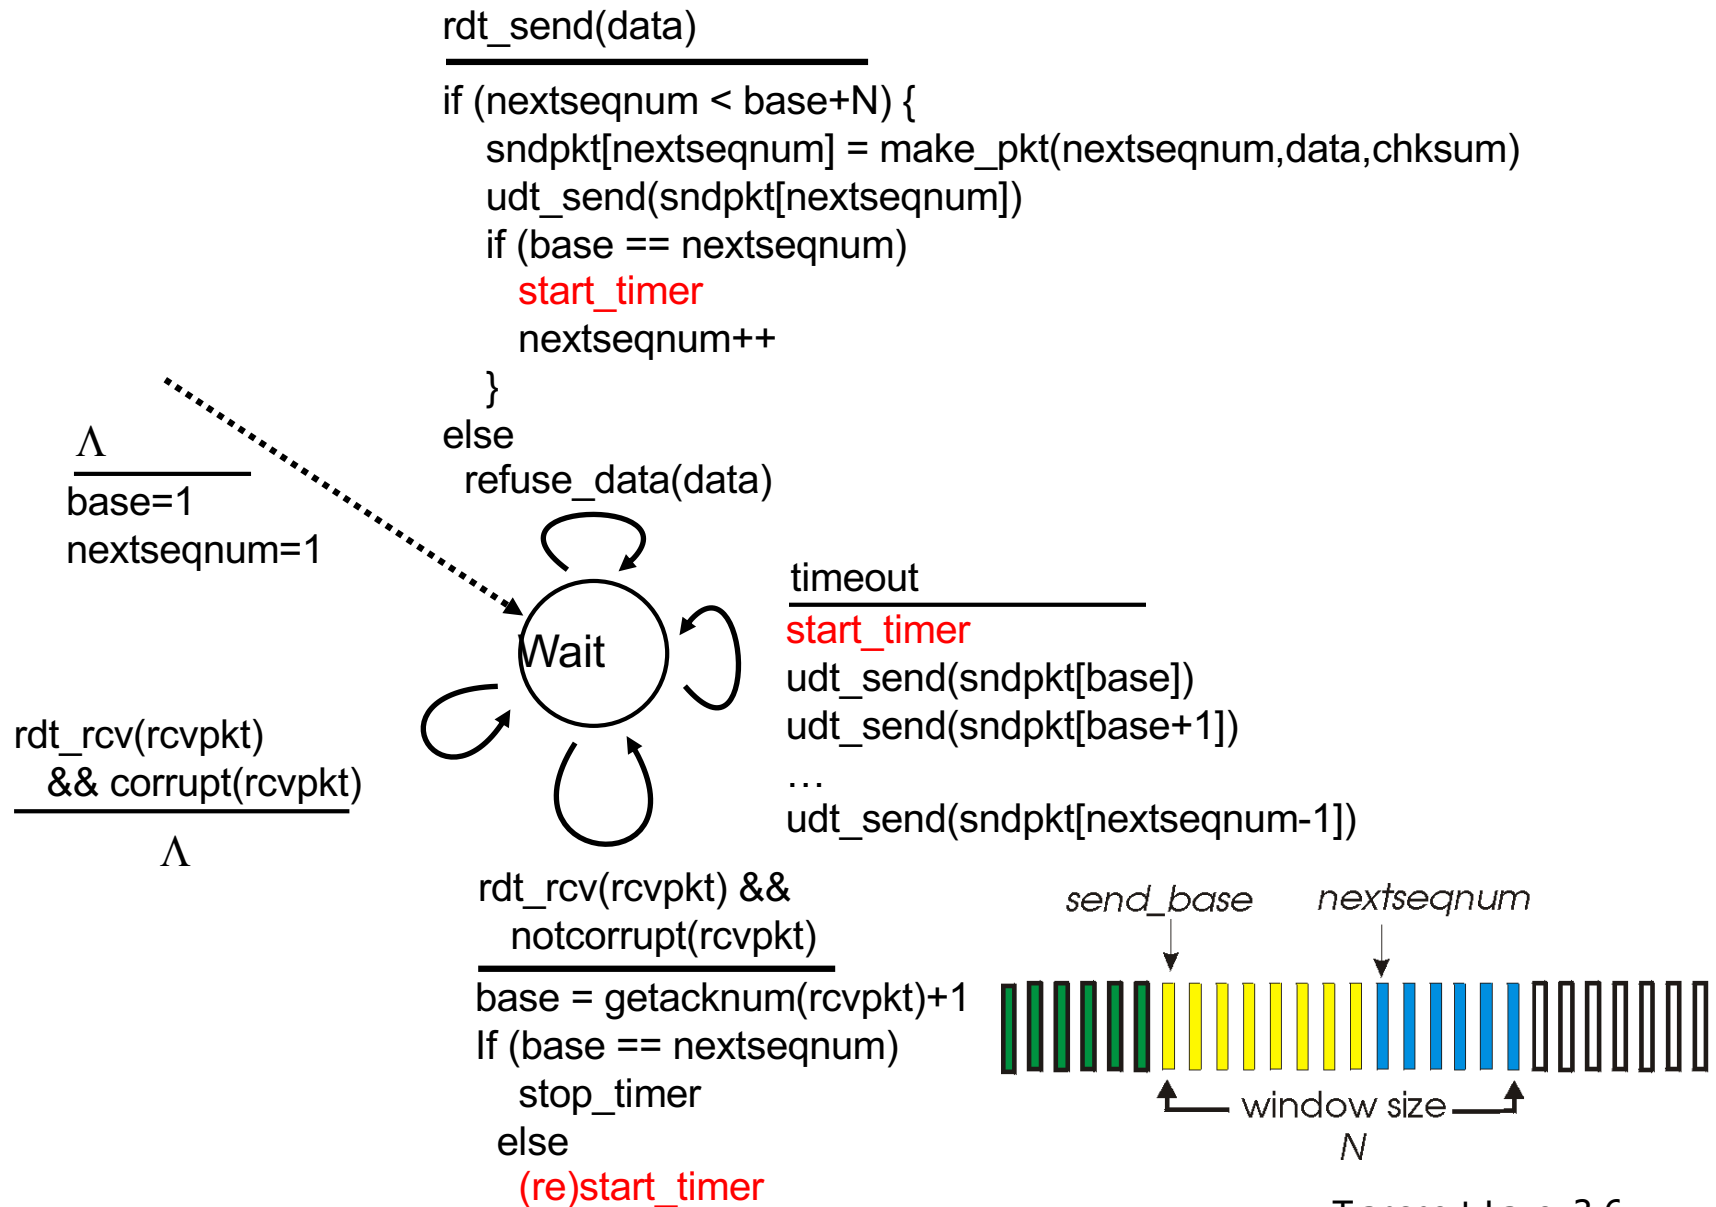
# Go-Back-N: sender

- k-bit seq # in pkt header
- “window” of up to N, consecutive unack’ed pkts allowed

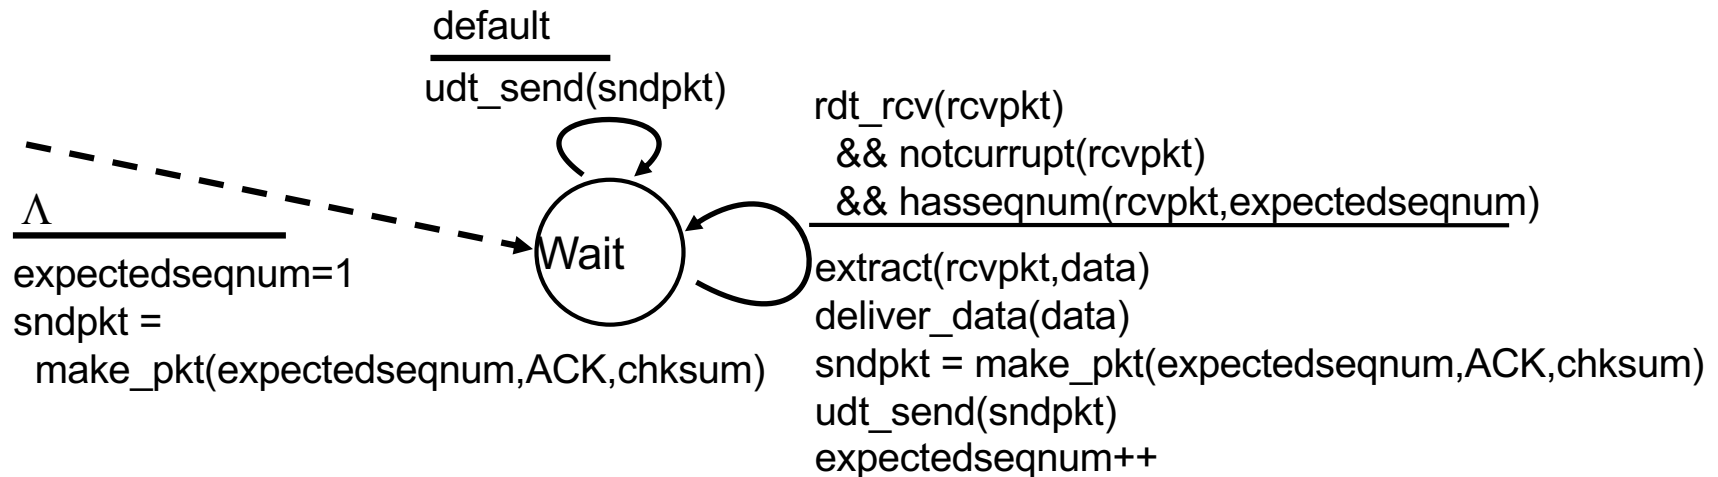


- ACK(n): ACKs all pkts up to, including seq # n - “cumulative ACK”
  - may receive duplicate ACKs (see receiver)
- timer for oldest in-flight pkt
- timeout(n): retransmit packet n and all higher seq # pkts in window

# GBN: sender extended FSM



# GBN: receiver extended FSM



ACK-only: always send ACK for correctly-received pkt with highest *in-order* seq #

- may generate duplicate ACKs
- need only remember **expectedseqnum**
- out-of-order pkt:
  - discard (don't buffer): *no receiver buffering!*
  - re-ACK pkt with highest in-order seq #

# GBN in action

sender window (N=4)

0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8

sender

send pkt0  
send pkt1  
send pkt2  
send pkt3  
(wait)

rcv ack0, send pkt4  
rcv ack1, send pkt5

ignore duplicate ACK



*pkt 2 timeout*

send pkt2  
send pkt3  
send pkt4  
send pkt5

receiver

receive pkt0, send ack0  
receive pkt1, send ack1

receive pkt3, discard,  
(re)send ack1

receive pkt4, discard,  
(re)send ack1  
receive pkt5, discard,  
(re)send ack1

rcv pkt2, deliver, send ack2  
rcv pkt3, deliver, send ack3  
rcv pkt4, deliver, send ack4  
rcv pkt5, deliver, send ack5



# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

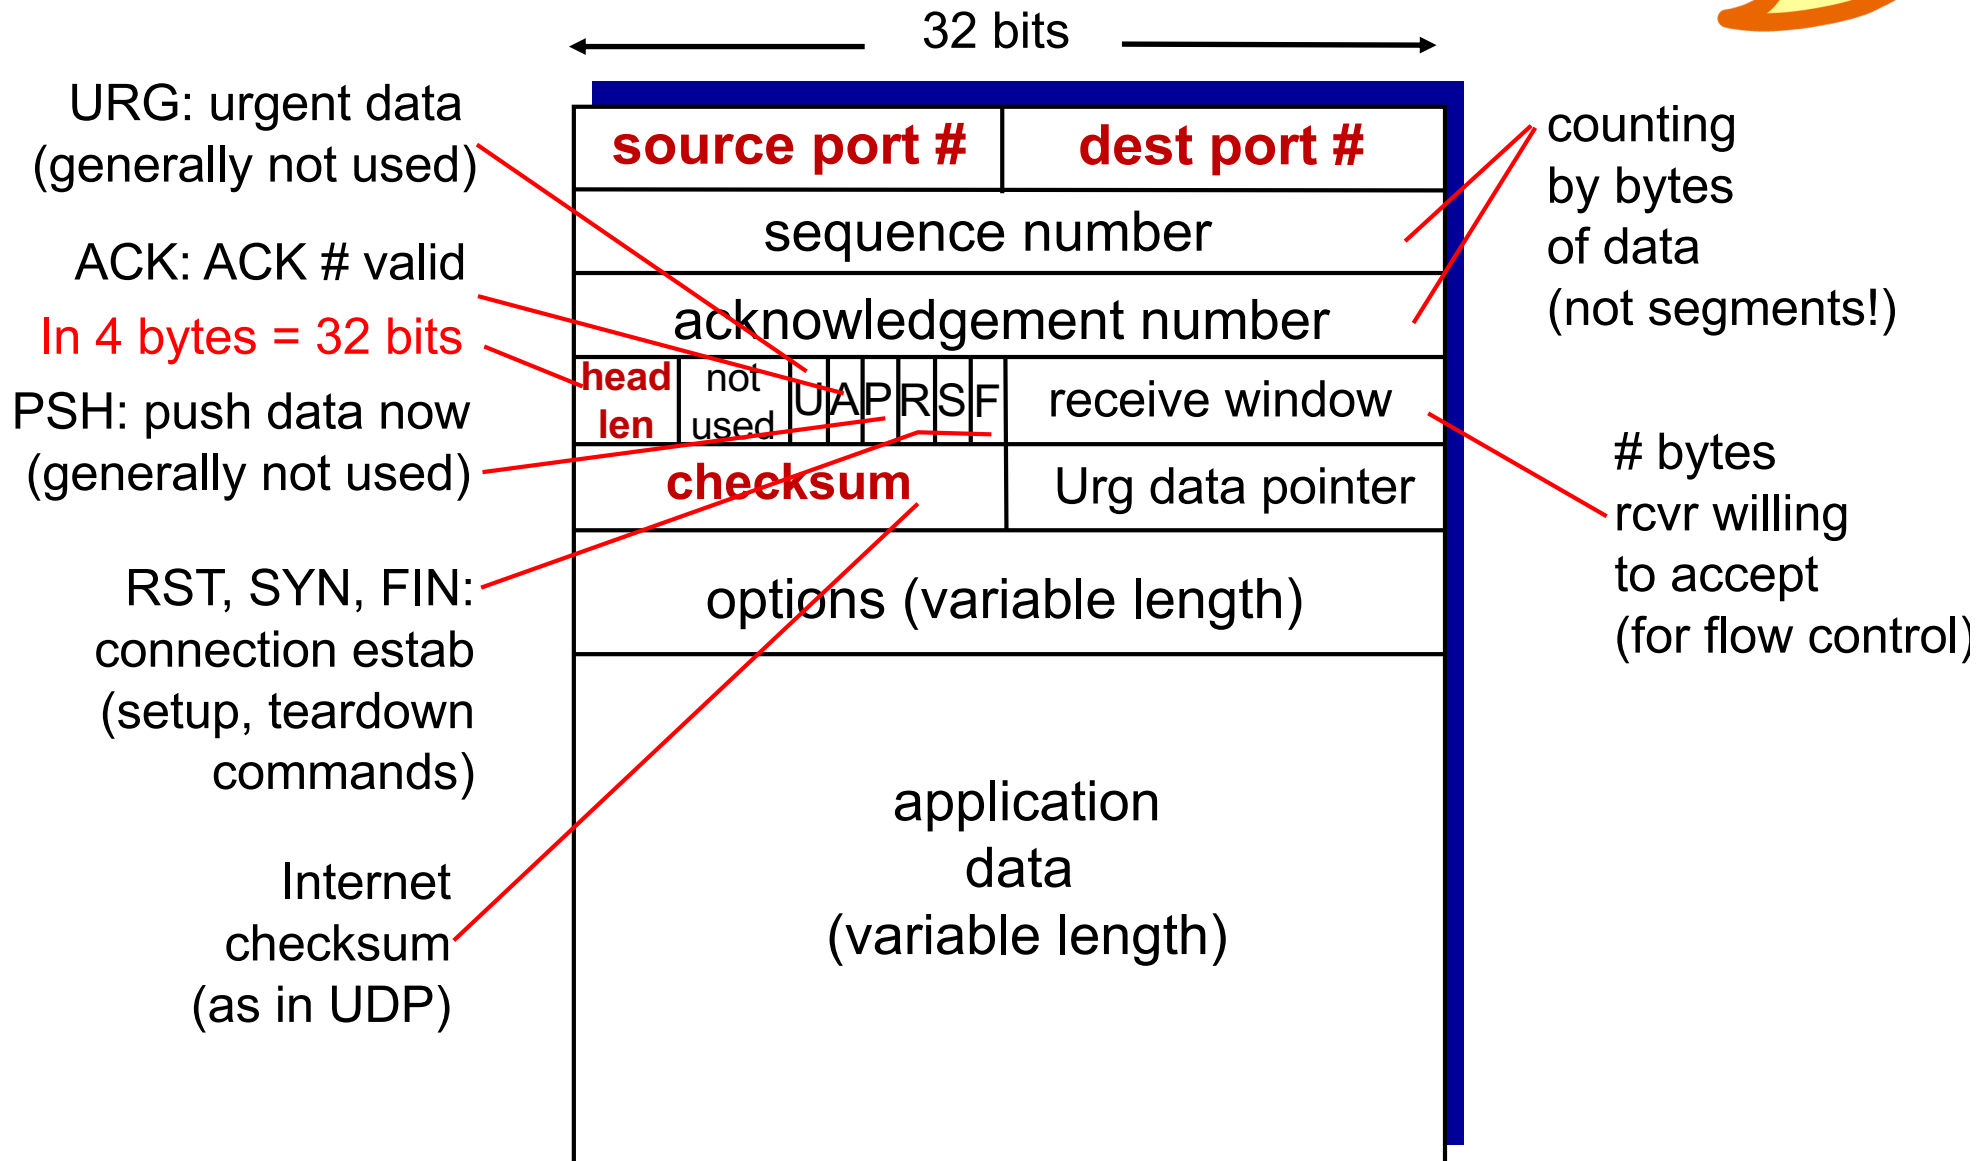
- segment structure
  - Segment Format
  - Sequence #
  - 3 way handshake
  - Timeout value estimation

# TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

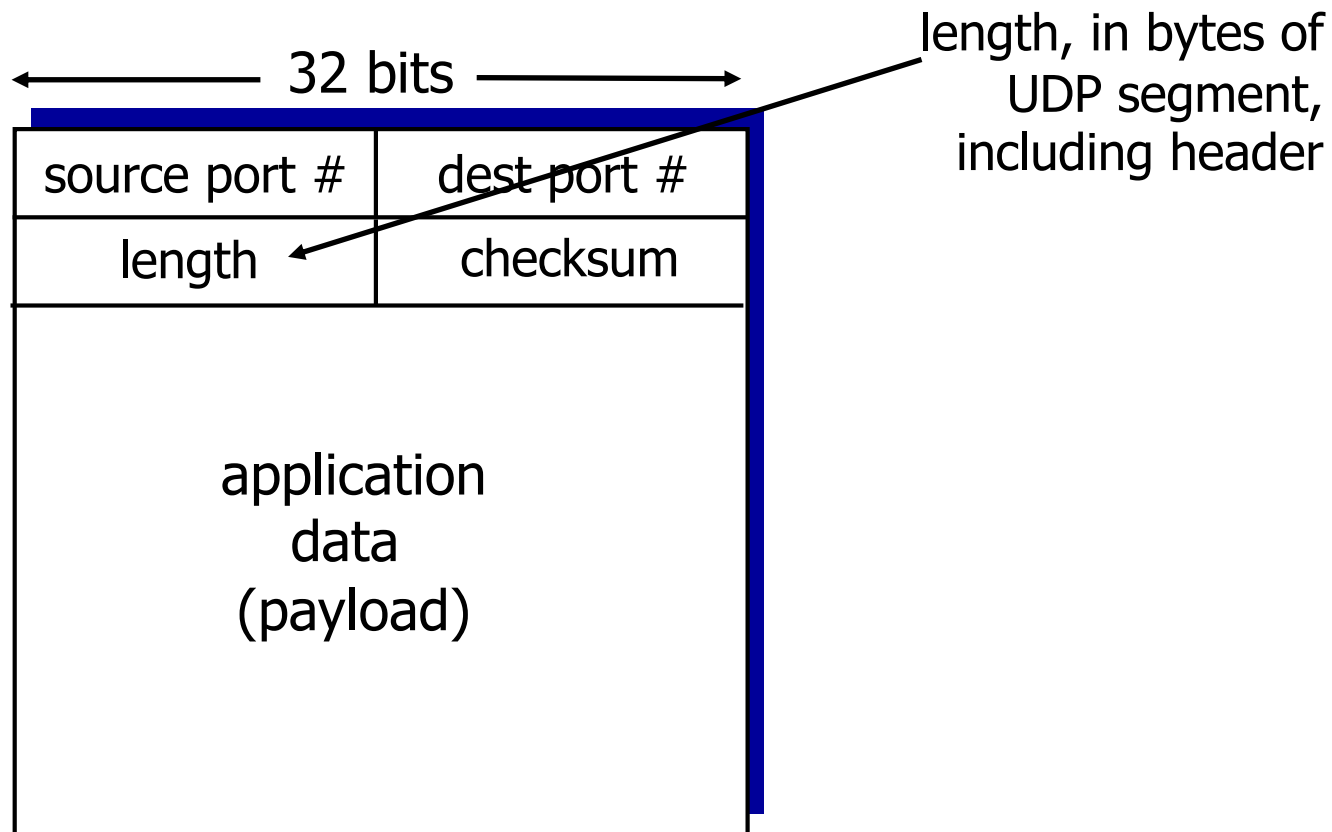
- ❖ **point-to-point:**
  - one sender, one receiver
- ❖ **reliable, in-order *byte stream*:**
  - no “message boundaries”
- ❖ **pipelined:**
  - TCP window size
- ❖ **full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size (e.g., 1460B)
  - MTU: layer 3 maximum transmission unit (e.g., 1500B for Ethernet)
- ❖ **connection-oriented:**
  - handshaking (exchange of control msgs) initializes sender, receiver state before data exchange

# TCP segment structure (20+ bytes)



Length of UDP header?

# UDP: segment header (8 bytes)



UDP segment format

# TCP seq. numbers, ACKs

## sequence numbers:

- byte stream “number” of first byte in segment’s data

## acknowledgements:

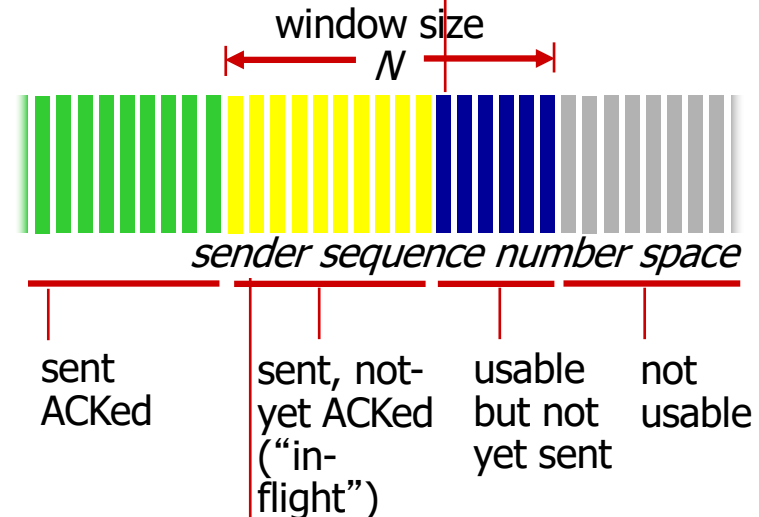
- seq # of next byte expected from other side
- cumulative ACK

**Q:** how receiver handles out-of-order segments

- A: TCP spec doesn’t say,  
- up to implementor
- rdt 3.0 & GBN & more

outgoing segment from sender

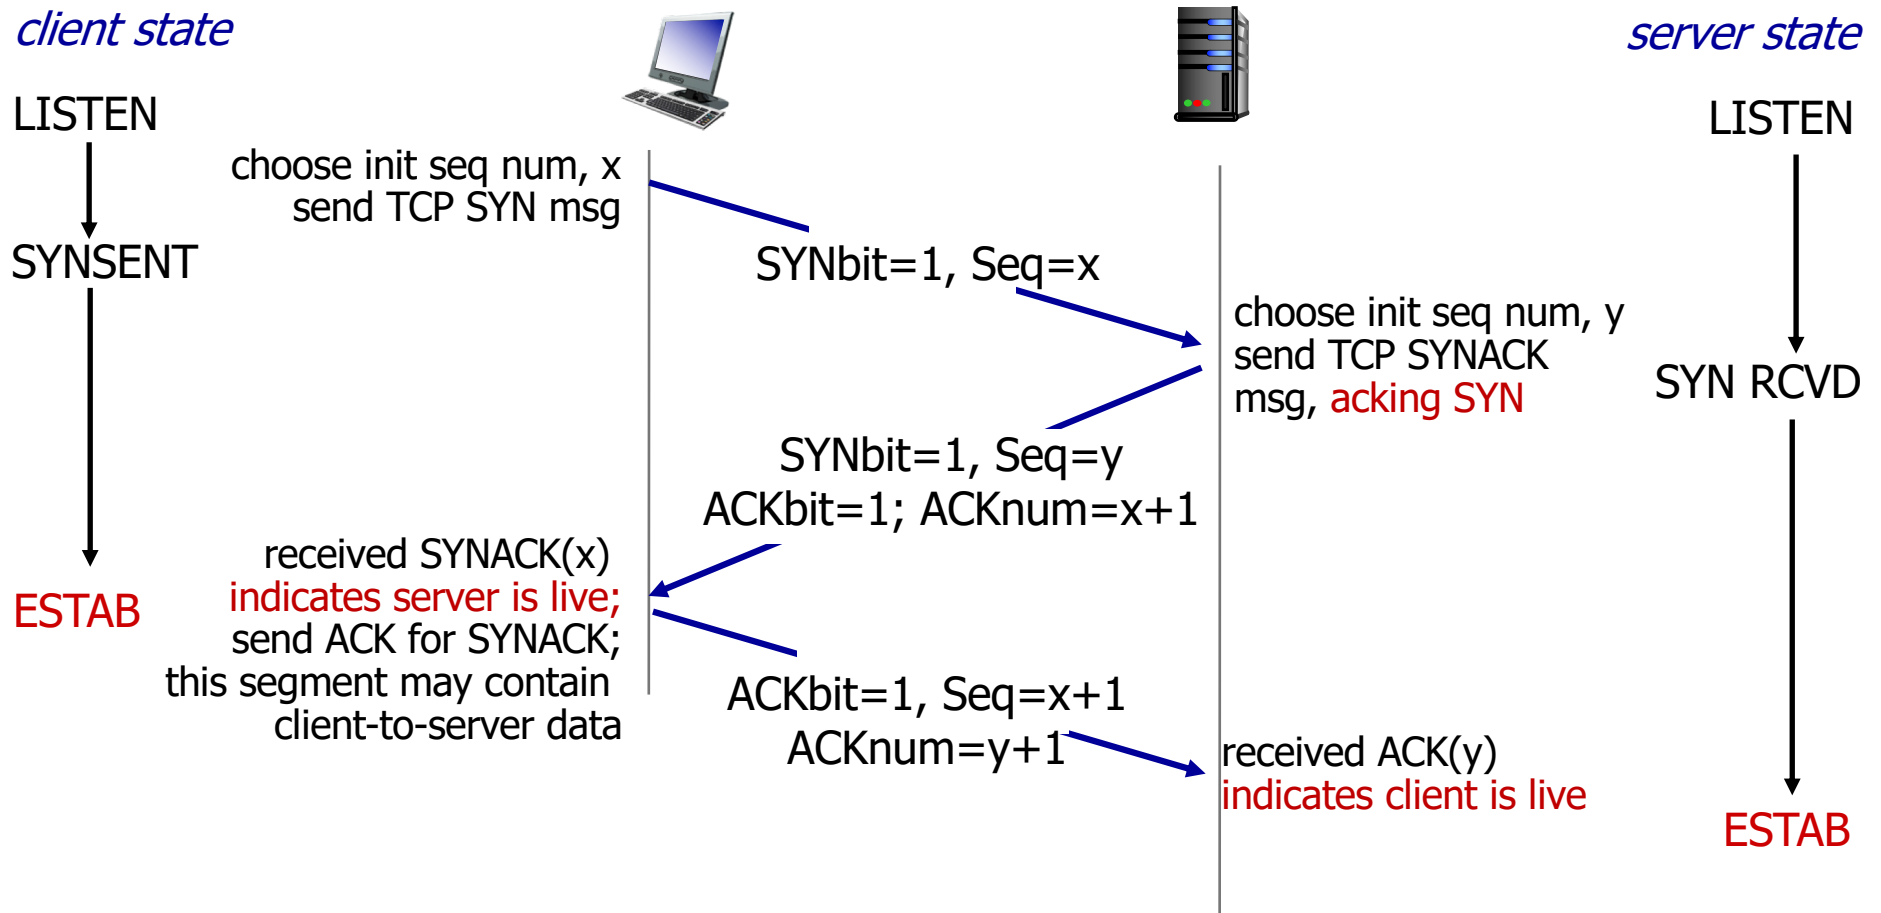
source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer

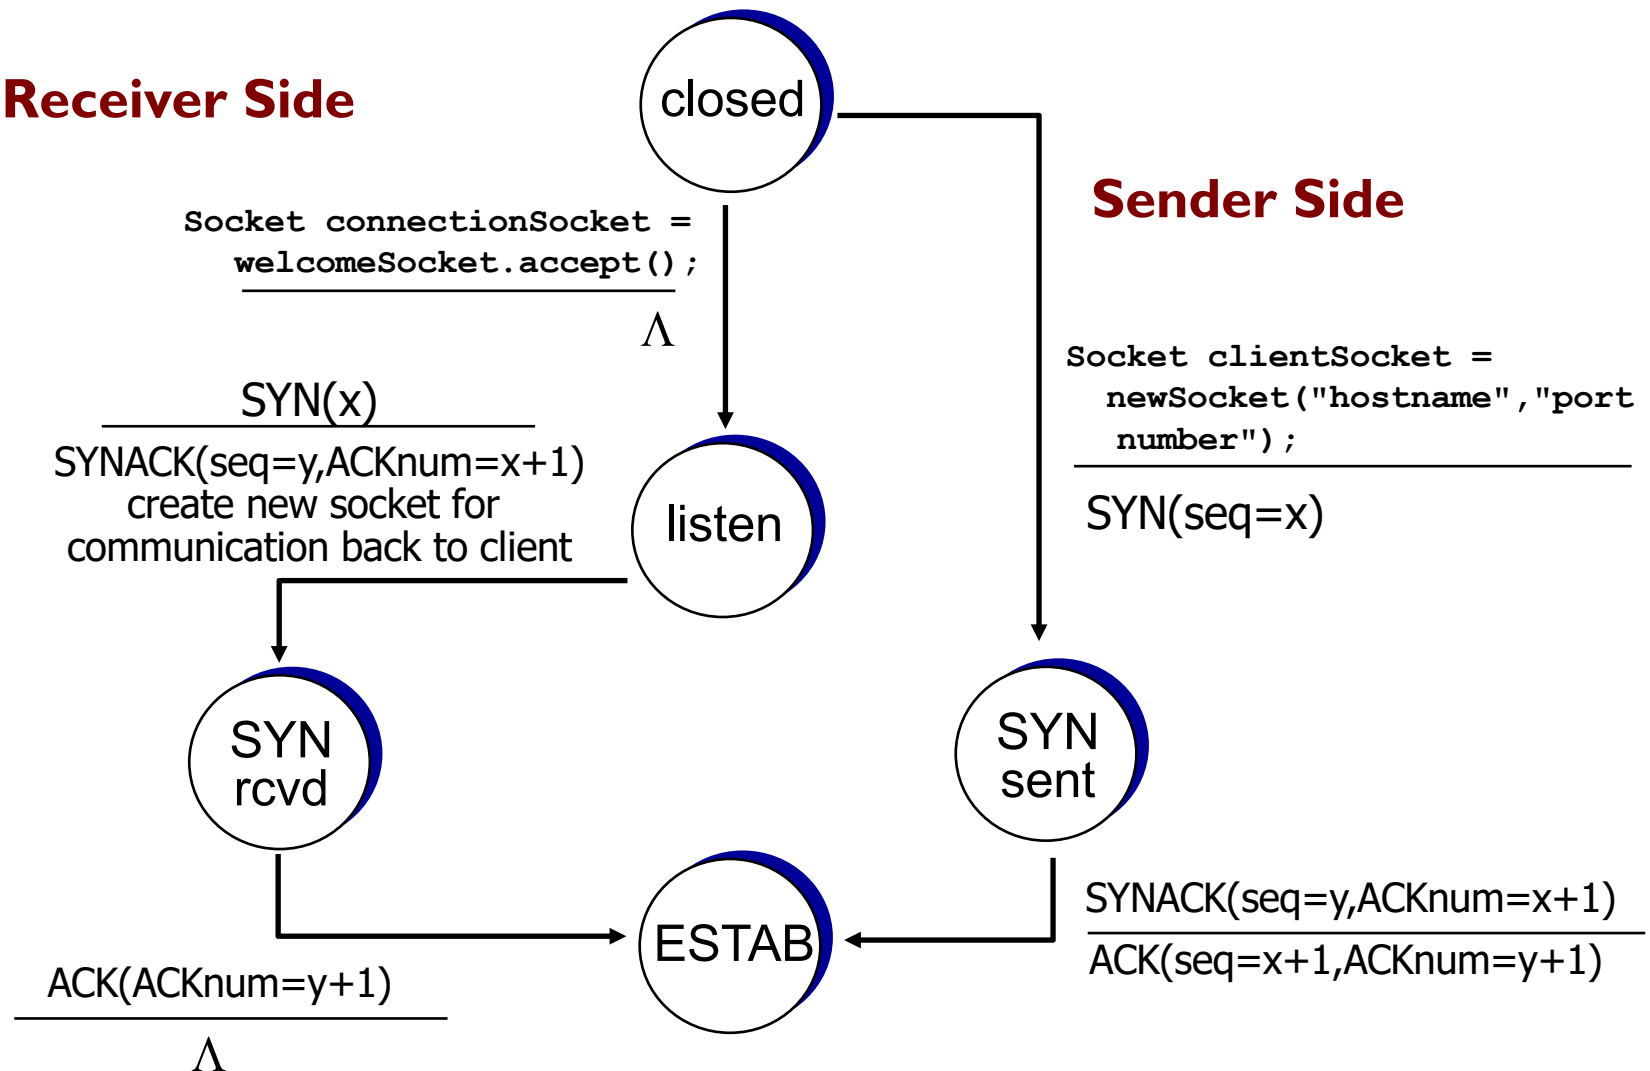
# TCP 3-way handshake



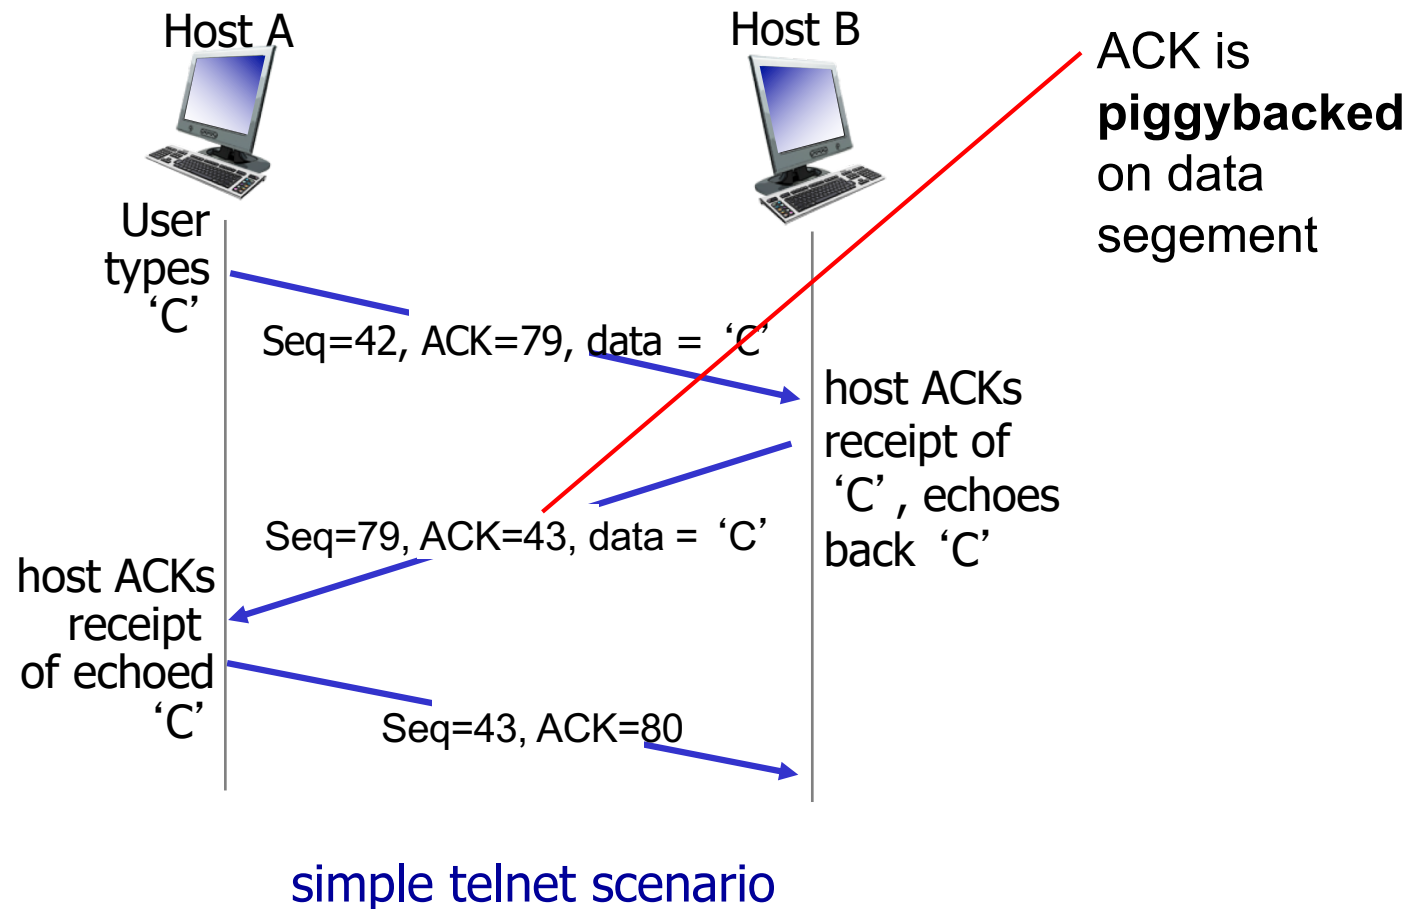
Random initial seq # are generated on both sides

# TCP 3-way handshake: FSM

## Receiver Side



# TCP seq. numbers, ACKs





# TCP round trip time, timeout



Q: how to set TCP timeout value?

- ❖ longer than RTT
  - but RTT varies
- ❖ *too short*: premature timeout, unnecessary retransmissions
- ❖ *too long*: slow reaction to segment loss

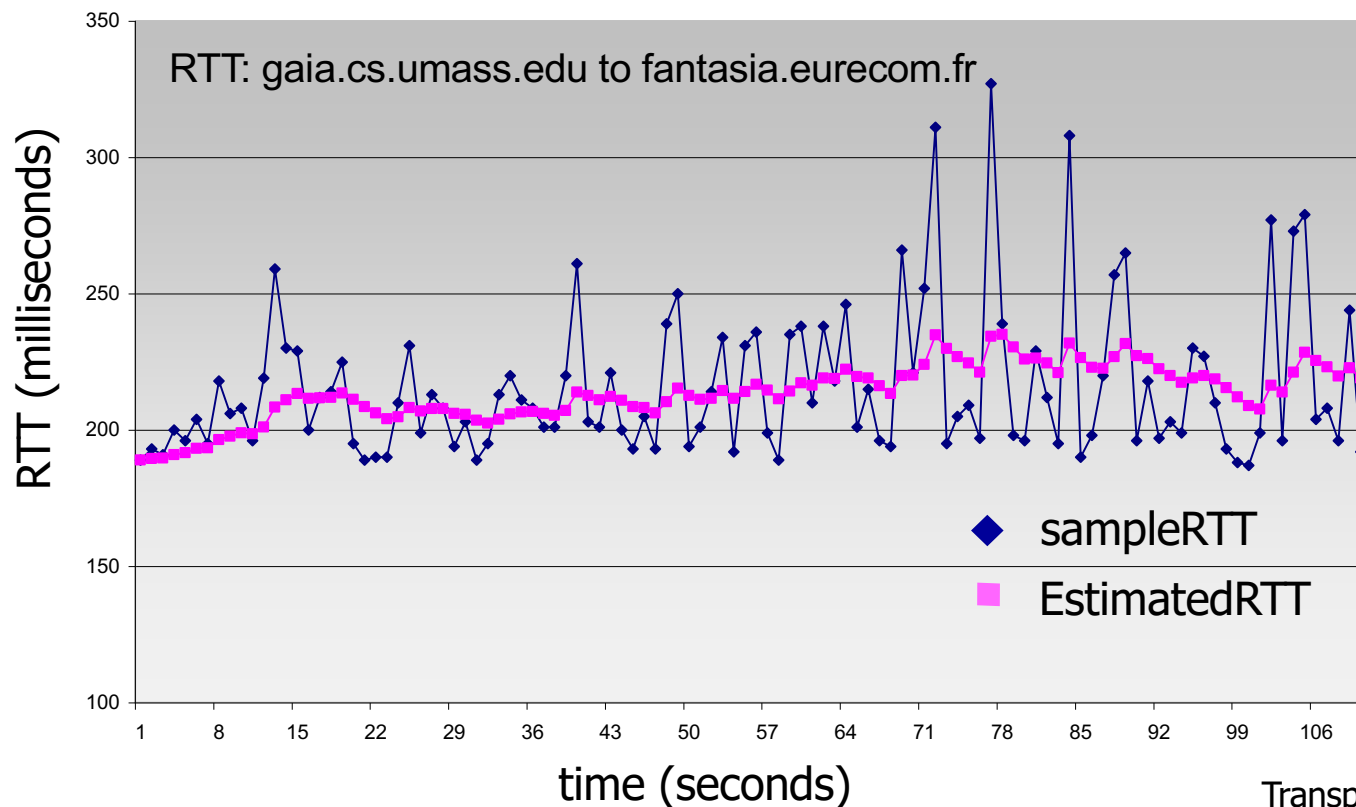
Q: how to estimate RTT?

- ❖ **SampleRTT**: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- ❖ **SampleRTT** will vary, want estimated RTT “smoother”
  - average several *recent* measurements, not just current **SampleRTT**

# TCP round trip time, timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- ❖ exponential weighted moving average
- ❖ influence of past sample decreases exponentially fast
- ❖ typical value:  $\alpha = 0.125$



# TCP round trip time, timeout

- ❖ **timeout interval:** **EstimatedRTT** plus “safety margin”
  - large variation in **EstimatedRTT** -> larger safety margin

- ❖ estimate SampleRTT deviation from EstimatedRTT:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$



↑  
estimated RTT

↑  
“safety margin”

# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
  - Segment Format
  - Sequence #
  - 3 way handshake
  - Timeout value estimation
- reliable data transfer
  - FSM
  - TCP Retransmissions

# TCP reliable data transfer (Sim. GBN)

- ❖ TCP creates rdt service on top of IP's unreliable service

- pipelined segments
- cumulative acks
- single retransmission timer

- ❖ retransmissions triggered by:

- timeout events
- duplicate acks

let's initially consider simplified TCP sender:

- ignore duplicate acks

# TCP sender events:

## *data rcvd from app:*

- ❖ create segment with seq #
- ❖ seq # is **byte-stream number** of first data byte in segment
- ❖ start timer if not already running
  - think of timer as for oldest unacked segment
  - expiration interval: `TimeoutInterval`

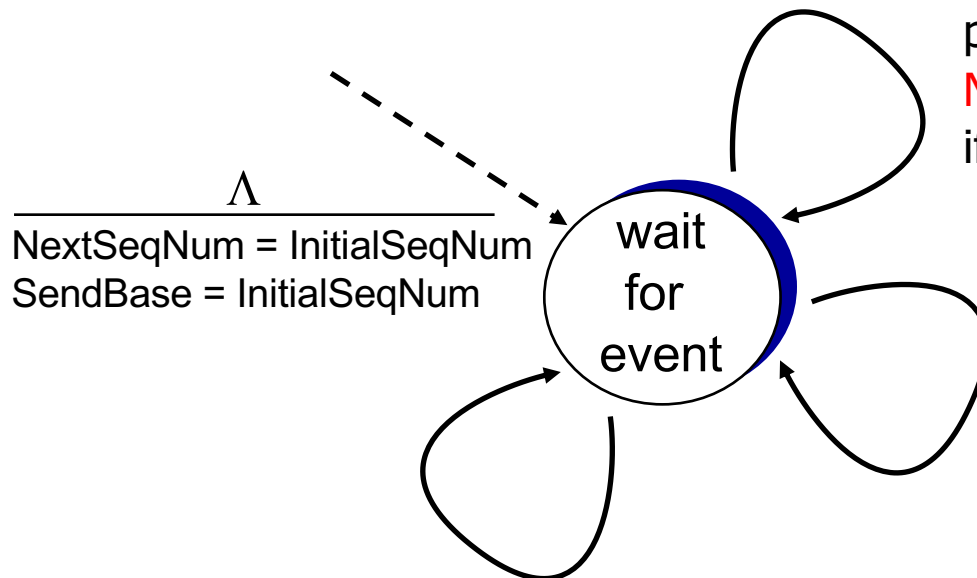
## *timeout:*

- ❖ retransmit segment that caused timeout
- ❖ restart timer

## *ack rcvd:*

- ❖ if ack acknowledges previously unacked segments
  - update what is known to be ACKed
  - start timer if there are still unacked segments

# TCP sender (simplified)



data received from application above

create segment, seq. #: NextSeqNum

pass segment to IP (i.e., “send”)

$\text{NextSeqNum} = \text{NextSeqNum} + \text{length}(\text{data})$

if (timer currently not running)

start timer

timeout

retransmit not-yet-acked segment

with smallest seq. #

start timer

ACK received, with ACK field value y

if (y > SendBase) {

SendBase = y

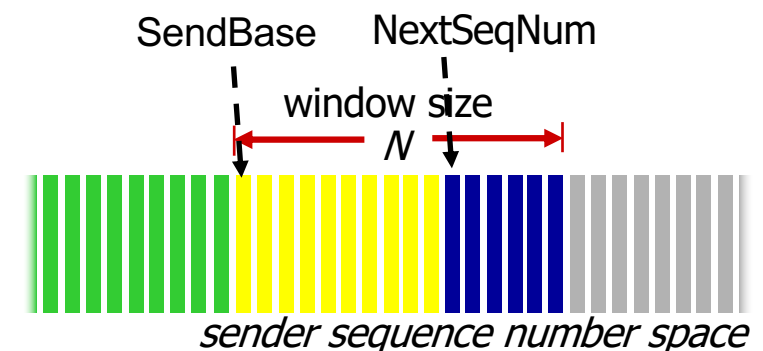
/\* SendBase-1: last cumulatively ACKed byte \*/

if (there are currently not-yet-acked segments)

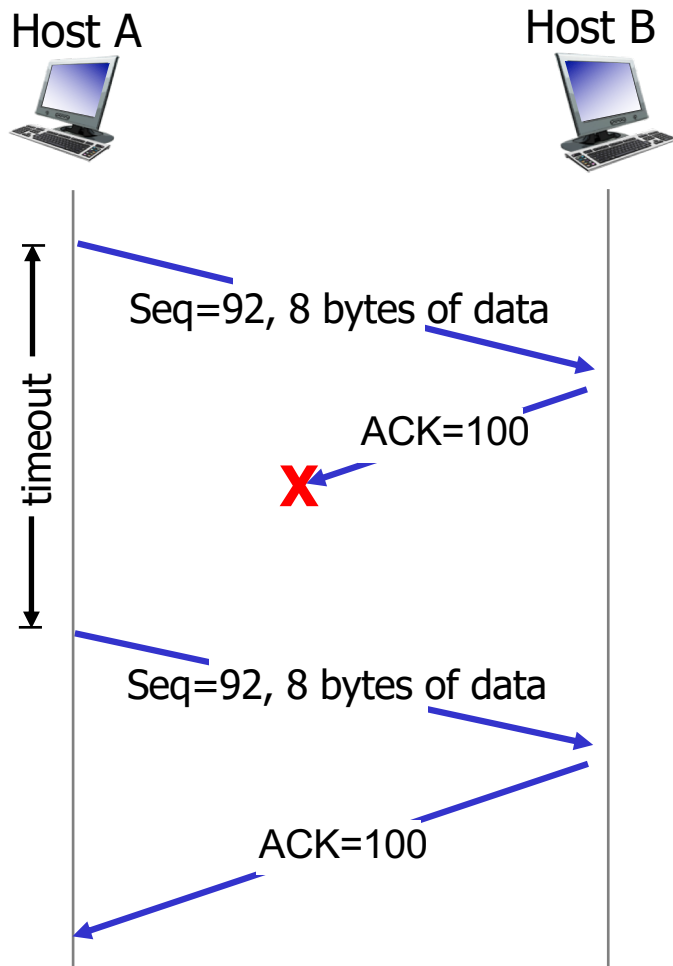
start timer

else stop timer

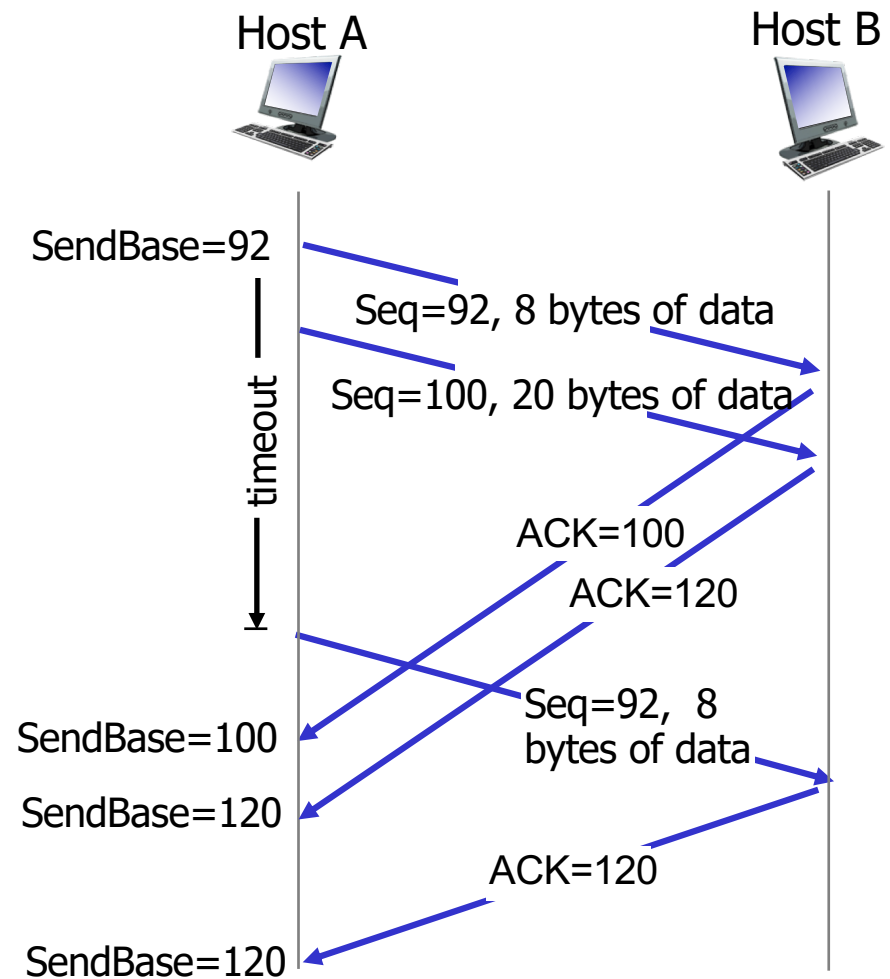
}



# TCP: retransmission scenarios



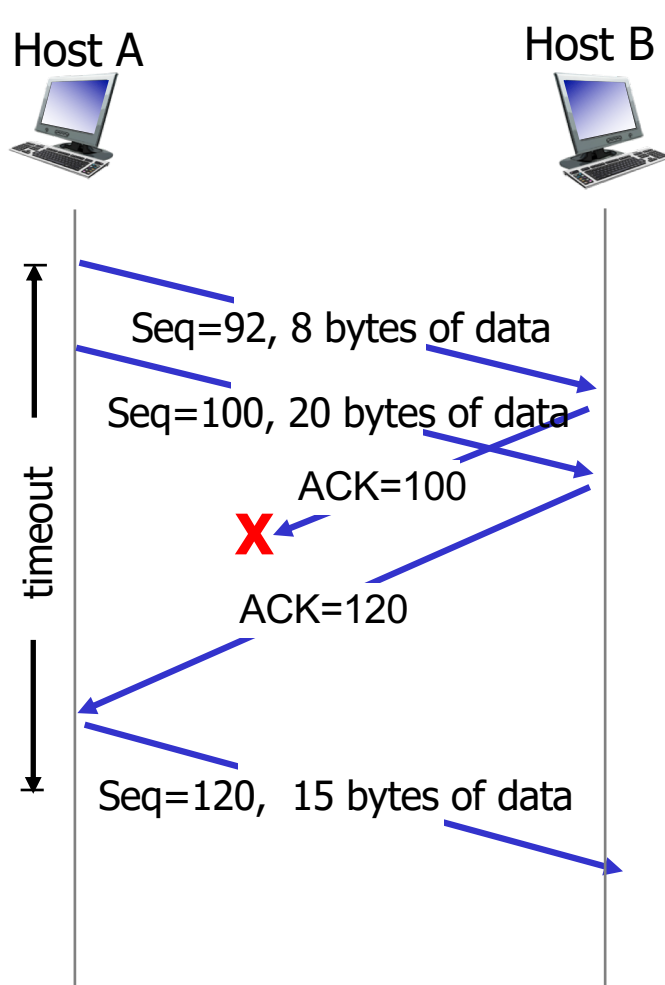
lost ACK scenario



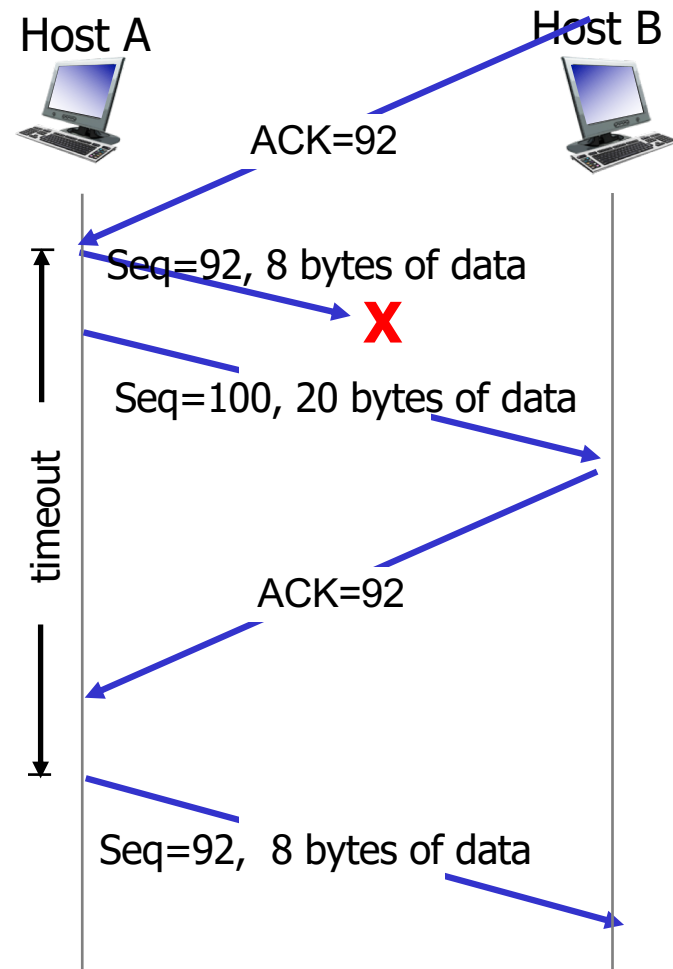
premature timeout



# TCP: retransmission scenarios



cumulative ACK



duplicate ACK

# TCP ACK generation [RFC 1122, RFC 2581]

<i>event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

# TCP fast retransmit

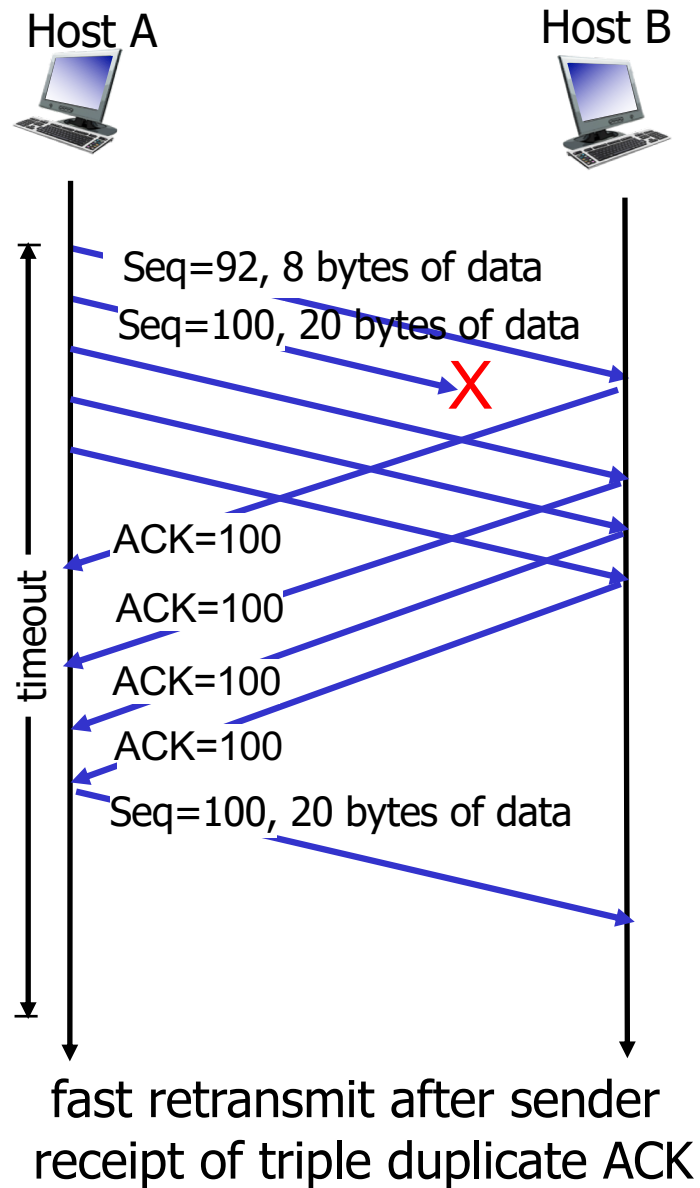
- ❖ time-out period often relatively long:
  - long delay before resending lost packet
- ❖ detect lost segments via duplicate ACKs.
  - sender often sends many segments back-to-back
  - if segment is lost, there will likely be many duplicate ACKs.

## *TCP fast retransmit*

if sender receives 3 dup ACKs for same data (“triple duplicate ACKs”), resend unacked segment with smallest seq #

- likely that unacked segment lost, so don't wait for timeout

# TCP fast retransmit



# Chapter 3: summary

- ❖ principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
- ❖ instantiation, implementation in the Internet
  - UDP
  - TCP

## next:

- ❖ leaving the network “edge” (application, transport layers)
- ❖ into the network “core”

# Questions?